

SQL in Quarto: three ways

USCOTS 2025 breakout session

Nicholas Horton (nhorton@amherst.edu) and Jo Hardin (jo.hardin@pomona.edu)

July 18, 2025

Today's example

The airlines database

Consider a database of all US flights between 2013 and 2015. The flight information was collected from the Bureau of Transportation Statistics, US Department of Transportation.

The database is a superset of the data in the `nycflights13` R package that tracks only flights departing airports serving New York City in 2013.



SQL connection

To set up a **SQL** connection, you need the location of the server (**host**) as well as a **username** and **password**. For example, you may want to use the subset of data from 2013 to 2015 which exists in a **SQL** database hosted by Ben Baumer to connect to the text *Modern Data Science in R*.

```

```{r}
con_mysql <- DBI::dbConnect(
 RMariaDB::MariaDB(),
 dbname = "airlines",
 host = Sys.getenv("MDSR_HOST"),
 user = Sys.getenv("MDSR_USER"),
 password = Sys.getenv("MDSR_PWD")
)
```

```

Using SQL in Quarto

We can write **SQL** code in three distinct ways:

1. Using the **DBI** package, we can send **SQL** queries through an **r** chunk (method used in the introductory presentation).
2. Using the package **dbplyr**, **R** will directly translate **dplyr** code into **SQL**.
3. Using a **sql** chunk, we can write actual **SQL** code inside a Quarto document.

1. SQL queries via the DBI package

- Look at the first few rows of the `flights` data.

```

```{r}
DBI::dbGetQuery(con_mysql,
 "SELECT * FROM flights LIMIT 8;")
```

```

| | year | month | day | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time | |
|---|-----------|---------|---------|----------|----------------|-----------|----------|----------------|-----------|
| 1 | 2013 | 10 | 1 | 2 | 10 | -8 | 453 | 505 | |
| 2 | 2013 | 10 | 1 | 4 | 2359 | 5 | 730 | 729 | |
| 3 | 2013 | 10 | 1 | 11 | 15 | -4 | 528 | 530 | |
| 4 | 2013 | 10 | 1 | 14 | 2355 | 19 | 544 | 540 | |
| 5 | 2013 | 10 | 1 | 16 | 17 | -1 | 515 | 525 | |
| 6 | 2013 | 10 | 1 | 22 | 20 | 2 | 552 | 554 | |
| 7 | 2013 | 10 | 1 | 29 | 35 | -6 | 808 | 816 | |
| 8 | 2013 | 10 | 1 | 29 | 35 | -6 | 449 | 458 | |
| | arr_delay | carrier | tailnum | flight | origin | dest | air_time | distance | cancelled |
| 1 | -12 | AA | N201AA | 2400 | LAX | DFW | 149 | 1235 | 0 |
| 2 | 1 | FL | N344AT | 710 | SFO | ATL | 247 | 2139 | 0 |

| | | | | | | | | | |
|---|-----|----|--------|------|-----|-----|-----|------|---|
| 3 | -2 | AA | N3KMAA | 1052 | SFO | DFW | 182 | 1464 | 0 |
| 4 | 4 | AA | N3ENAA | 2392 | SEA | ORD | 191 | 1721 | 0 |
| 5 | -10 | UA | N38473 | 1614 | LAX | IAH | 157 | 1379 | 0 |
| 6 | -2 | UA | N458UA | 291 | SFO | IAH | 188 | 1635 | 0 |
| 7 | -8 | US | N551UW | 436 | LAX | CLT | 256 | 2125 | 0 |
| 8 | -9 | AS | N402AS | 108 | ANC | SEA | 181 | 1448 | 0 |

| | diverted | hour | minute | time_hour |
|---|----------|------|--------|---------------------|
| 1 | 0 | 0 | 10 | 2013-10-01 00:10:00 |
| 2 | 0 | 23 | 59 | 2013-10-01 23:59:00 |
| 3 | 0 | 0 | 15 | 2013-10-01 00:15:00 |
| 4 | 0 | 23 | 55 | 2013-10-01 23:55:00 |
| 5 | 0 | 0 | 17 | 2013-10-01 00:17:00 |
| 6 | 0 | 0 | 20 | 2013-10-01 00:20:00 |
| 7 | 0 | 0 | 35 | 2013-10-01 00:35:00 |
| 8 | 0 | 0 | 35 | 2013-10-01 00:35:00 |

1. SQL queries via the DBI package

- How many flights per year are in the `flights` table?

```
```{r}
#| cache: true

DBI::dbGetQuery(con_mysql,
 "SELECT year, COUNT(*) AS num_flights
 FROM flights
 GROUP BY year
 ORDER BY num_flights;")
```
```

```

  year num_flights
1 2015      5819079
2 2014      5819811
3 2013      6369482
```

1. SQL queries via the DBI package

- What is the average Departure Delay (and number of flights) for each destination?

```

```{r}
#| cache: true

DBI::dbGetQuery(con_mysql,
 "SELECT dest,
 AVG(dep_delay) AS mean_delay,
 COUNT(*) AS num_flights
 FROM flights
 GROUP BY dest
 LIMIT 8;")
```

```

| | dest | mean_delay | num_flights |
|---|------|------------|-------------|
| 1 | ABE | 9.9056 | 7253 |
| 2 | ABI | 8.9290 | 8114 |
| 3 | ABQ | 11.7639 | 74189 |
| 4 | ABR | 2.4078 | 2239 |
| 5 | ABY | 9.5748 | 3001 |
| 6 | ACK | 5.5521 | 1295 |
| 7 | ACT | 10.2526 | 5225 |
| 8 | ACV | 12.5367 | 7785 |

2. Translating dplyr code into SQL

```

flights_tbl <- dplyr::tbl(con_mysql, "flights")

flights_tbl |>
  head()

```

```

# Source:   SQL [?? x 21]
# Database: mysql [mdsr_public@mdsr.crcbo51tmesf.us-east-2.rds.amazonaws.com:3306/airlines]
  year month   day dep_time sched_dep_time dep_delay arr_time
  <int> <int> <int>   <int>         <int>         <int>   <int>
1  2013    10     1         2             10          -8       453
2  2013    10     1         4            2359           5       730
3  2013    10     1        11             15          -4       528
4  2013    10     1        14            2355          19       544
5  2013    10     1        16             17           -1       515
6  2013    10     1        22             20           2       552
# i 14 more variables: sched_arr_time <int>, arr_delay <int>,

```

```
#   carrier <chr>, tailnum <chr>, flight <int>, origin <chr>,
#   dest <chr>, air_time <int>, distance <int>, cancelled <int>,
#   diverted <int>, hour <int>, minute <int>, time_hour <dtm>
```

2. Translating dplyr code into SQL

- Over what years is the flights data taken?

```
yrs <- flights_tbl |>
  summarize(min_year = min(year, na.rm = TRUE),
            max_year = max(year, na.rm = TRUE))

yrs
```

```
# Source:   SQL [?? x 2]
# Database: mysql [mdsr_public@mdsr.crcbo51tmesf.us-east-2.rds.amazonaws.com:3306/airlines]
   min_year max_year
   <int>     <int>
1     2013     2015
```

2. Translating dplyr code into SQL

Because `flights_tbl` is not actually a `data.frame` in **R** (but instead a `tbl` in **SQL**), the work that was done above was actually performed in **SQL**. To see the **SQL** code, we can use the function `show_query`.

```
```{r}
dplyr::show_query(yrs)
```
```

```
<SQL>
SELECT MIN(`year`) AS `min_year`, MAX(`year`) AS `max_year`
FROM `flights`
```

2. Translating dplyr code into SQL

- Create a data set containing only flights between LAX and BOS in 2015.

```

```{r}
#| cache: true

la_bos <- flights_tbl |>
 filter(year == 2015 & ((origin == "LAX" & dest == "BOS") |
 (origin == "BOS" & dest == "LAX")))

dplyr::show_query(la_bos)
```

```

```

<SQL>
SELECT `flights`.*
FROM `flights`
WHERE (`year` = 2015.0 AND ((`origin` = 'LAX' AND `dest` = 'BOS') OR (`origin` = 'BOS' AND `

```

2. Translating dplyr code into SQL

- What is the average Departure Delay (and number of flights) for each destination?

```

```{r}
#| cache: true

aveDepDel <- flights_tbl |>
 group_by(dest) |>
 summarize(mean_delay = mean(dep_delay), num_flights = n())

aveDepDel

dplyr::show_query(aveDepDel)
```

```

```

# Source:   SQL [?? x 3]
# Database: mysql [mdsr_public@mdsr.crcbo51tmesf.us-east-2.rds.amazonaws.com:3306/airlines]
  dest mean_delay num_flights
  <chr>      <dbl>      <int64>
1 ABE         9.91         7253
2 ABI         8.93         8114
3 ABQ        11.8         74189
4 ABR         2.41         2239
5 ABY         9.57         3001
6 ACK         5.55         1295

```

```

7 ACT      10.3      5225
8 ACV      12.5      7785
9 ACY      12.0      4311
10 ADK     -1.14      313
# i more rows
<SQL>
SELECT `dest`, AVG(`dep_delay`) AS `mean_delay`, COUNT(*) AS `num_flights`
FROM `flights`
GROUP BY `dest`

```

2. Translating dplyr code into SQL

- **dbplyr** doesn't translate every **R** command into **SQL**.
- **SQL** is not a statistical software and doesn't, for example, have a mechanism for creating data visualizations.
- track which **R** commands are connected to **SQL** at the **dbplyr** reference sheet.

3. Direct SQL queries via sql chunks

SQL queries can be written directly inside a `sql` chunk in Quarto.

```

```{sql}
#| connection: con_mysql

SELECT * FROM flights LIMIT 8;
```

```

Table 1: 8 records

| year | month | day | dep_scheduled | dep_delayed | dep_delay | arr_delay | carrier | flight | origin | dest | aircraft | distance | cancelled | diverted | time_hour | | | |
|------|-------|-----|---------------|-------------|-----------|-----------|---------|--------|---------|------|----------|----------|-----------|----------|-----------|----|---------------------|---------------------|
| 2013 | 1 | 2 | 10 | -8 | 453 | 505 | - | AA | N20240 | LAX | DFW | 1235 | 0 | 0 | 0 | 10 | 2013-10-01 00:10:00 | |
| 2013 | 1 | 4 | 2359 | 5 | 730 | 729 | 1 | FL | N3447AT | SFO | ATL | 1247 | 2139 | 0 | 0 | 23 | 59 | 2013-10-01 23:59:00 |

| year | month | day | dep_sch | dep_time | dep_delay | dep_delay_min | arr_delay | carrier | tailnum | flight | origin | dest | air_time | distance | cancelled | diverted | minute | time_hour |
|------|-------|-----|---------|----------|-----------|---------------|-----------|---------|---------|----------|---------|------|----------|----------|-----------|----------|--------|---------------------|
| 2013 | 10 | 1 | 11 | 15 | -4 | 528 | 530 | -2 | AA | N3K10528 | FODFW | 82 | 1464 | 0 | 0 | 0 | 15 | 2013-10-01 00:15:00 |
| 2013 | 10 | 1 | 14 | 2355 | 19 | 544 | 540 | 4 | AA | N3E23925 | SEAORD | 91 | 1721 | 0 | 0 | 23 | 55 | 2013-10-01 23:55:00 |
| 2013 | 10 | 1 | 16 | 17 | -1 | 515 | 525 | - | UA | N3841614 | LAXAH | 57 | 1379 | 0 | 0 | 0 | 17 | 2013-10-01 00:17:00 |
| 2013 | 10 | 1 | 22 | 20 | 2 | 552 | 554 | -2 | UA | N45291 | ASFOIAH | 88 | 1635 | 0 | 0 | 0 | 20 | 2013-10-01 00:20:00 |
| 2013 | 10 | 1 | 29 | 35 | -6 | 808 | 816 | -8 | US | N55436 | LAXCL | 256 | 2125 | 0 | 0 | 0 | 35 | 2013-10-01 00:35:00 |
| 2013 | 10 | 1 | 29 | 35 | -6 | 449 | 458 | -9 | AS | N40218 | ASANCEA | 81 | 1448 | 0 | 0 | 0 | 35 | 2013-10-01 00:35:00 |

3. Direct SQL queries via sql chunks

SQL queries can be written directly inside a sql chunk in Quarto.

```
```{sql}
#| connection: con_mysql
#| cache: true

SELECT year, count(*) AS num_flights
 FROM flights
 GROUP BY year
 ORDER BY num_flights;
```
```


Table 2: 3 records

| year | num_flights |
|------|-------------|
| 2015 | 5819079 |
| 2014 | 5819811 |
| 2013 | 6369482 |

3. Direct SQL queries via sql chunks

- What is the average Departure Delay (and number of flights) for each destination?

```

```{sql}
#| connection: con_mysql
#| cache: true

SELECT dest,
 AVG(dep_delay) AS mean_delay,
 COUNT(*) AS num_flights
FROM flights
GROUP BY dest
LIMIT 8;
```

```

Table 3: 8 records

| dest | mean_delay | num_flights |
|------|------------|-------------|
| ABE | 9.9056 | 7253 |
| ABI | 8.9290 | 8114 |
| ABQ | 11.7639 | 74189 |
| ABR | 2.4078 | 2239 |
| ABY | 9.5748 | 3001 |
| ACK | 5.5521 | 1295 |
| ACT | 10.2526 | 5225 |
| ACV | 12.5367 | 7785 |

Good practice

Always a good idea to terminate the **SQL** connection when you are done with it.

```
RMariaDB::dbDisconnect(con_mysql, shutdown = TRUE)
```

Parquet

- Connecting to the local parquet files is very similar to connecting to a remote **SQL** database.
- Note that the datasets themselves have slightly different variable names (not ideal for teaching!!!)

Connecting to the parquet files via DuckDB

Start an in-memory database using DuckDB. The function `duckdb()` comes from the **duckdb** package.

```
con_parq <- DBI::dbConnect(duckdb())
```

1. SQL queries via the DBI package

```
```{r}
DBI::dbGetQuery(con_parq,
 "SELECT *
 FROM read_parquet('activity/data_airlines/Year*/*.parquet')
 LIMIT 8;")
```
```

| | Year | Quarter | Month | DayofMonth | DayOfWeek | FlightDate | Reporting_Airline |
|---|------|---------|-------|------------|-----------|------------|-------------------|
| 1 | 2023 | 1 | 1 | 2 | 1 | 2023-01-02 | 9E |
| 2 | 2023 | 1 | 1 | 3 | 2 | 2023-01-03 | 9E |
| 3 | 2023 | 1 | 1 | 4 | 3 | 2023-01-04 | 9E |
| 4 | 2023 | 1 | 1 | 5 | 4 | 2023-01-05 | 9E |
| 5 | 2023 | 1 | 1 | 6 | 5 | 2023-01-06 | 9E |
| 6 | 2023 | 1 | 1 | 7 | 6 | 2023-01-07 | 9E |
| 7 | 2023 | 1 | 1 | 14 | 6 | 2023-01-14 | 9E |
| 8 | 2023 | 1 | 1 | 21 | 6 | 2023-01-21 | 9E |

| | DOT_ID_Reporting_Airline | IATA_CODE_Reporting_Airline | Tail_Number |
|---|--------------------------|-----------------------------|-------------|
| 1 | | 20363 | 9E N605LR |
| 2 | | 20363 | 9E N605LR |

| | | | |
|---|-------|----|--------|
| 3 | 20363 | 9E | N331PQ |
| 4 | 20363 | 9E | N906XJ |
| 5 | 20363 | 9E | N337PQ |
| 6 | 20363 | 9E | N336PQ |
| 7 | 20363 | 9E | N311PQ |
| 8 | 20363 | 9E | N917XJ |

| Flight_Number_Reporting_Airline | OriginAirportID | OriginAirportSeqID |
|---------------------------------|-----------------|--------------------|
| 1 | 4628 | 10529 |
| 2 | 4628 | 10529 |
| 3 | 4628 | 10529 |
| 4 | 4628 | 10529 |
| 5 | 4628 | 10529 |
| 6 | 4628 | 10529 |
| 7 | 4628 | 12953 |
| 8 | 4628 | 12953 |

| OriginCityMarketID | Origin | OriginCityName | OriginState | OriginStateFips |
|--------------------|--------|----------------|--------------|-----------------|
| 1 | 30529 | BDL | Hartford, CT | CT |
| 2 | 30529 | BDL | Hartford, CT | CT |
| 3 | 30529 | BDL | Hartford, CT | CT |
| 4 | 30529 | BDL | Hartford, CT | CT |
| 5 | 30529 | BDL | Hartford, CT | CT |
| 6 | 30529 | BDL | Hartford, CT | CT |
| 7 | 31703 | LGA | New York, NY | NY |
| 8 | 31703 | LGA | New York, NY | NY |

| OriginStateName | OriginWac | DestAirportID | DestAirportSeqID | DestCityMarketID |
|-----------------|-------------|---------------|------------------|------------------|
| 1 | Connecticut | 11 | 12953 | 1295304 |
| 2 | Connecticut | 11 | 12953 | 1295304 |
| 3 | Connecticut | 11 | 12953 | 1295304 |
| 4 | Connecticut | 11 | 12953 | 1295304 |
| 5 | Connecticut | 11 | 12953 | 1295304 |
| 6 | Connecticut | 11 | 12953 | 1295304 |
| 7 | New York | 22 | 11193 | 1119302 |
| 8 | New York | 22 | 11193 | 1119302 |

| Dest | DestCityName | DestState | DestStateFips | DestStateName | DestWac | CRSDepTime |
|------|--------------|----------------|---------------|---------------|----------|------------|
| 1 | LGA | New York, NY | NY | 36 | New York | 22 |
| 2 | LGA | New York, NY | NY | 36 | New York | 22 |
| 3 | LGA | New York, NY | NY | 36 | New York | 22 |
| 4 | LGA | New York, NY | NY | 36 | New York | 22 |
| 5 | LGA | New York, NY | NY | 36 | New York | 22 |
| 6 | LGA | New York, NY | NY | 36 | New York | 22 |
| 7 | CVG | Cincinnati, OH | KY | 21 | Kentucky | 52 |
| 8 | CVG | Cincinnati, OH | KY | 21 | Kentucky | 52 |

| DepTime | DepDelay | DepDelayMinutes | DepDel15 | DepartureDelayGroups | DepTimeBlk |
|---------|----------|-----------------|----------|----------------------|------------|
|---------|----------|-----------------|----------|----------------------|------------|

| | | | | | | |
|---|------|-----|---|---|----|-----------|
| 1 | 0757 | -3 | 0 | 0 | -1 | 0800-0859 |
| 2 | 0755 | -5 | 0 | 0 | -1 | 0800-0859 |
| 3 | 0755 | -5 | 0 | 0 | -1 | 0800-0859 |
| 4 | 0754 | -6 | 0 | 0 | -1 | 0800-0859 |
| 5 | 0759 | -1 | 0 | 0 | -1 | 0800-0859 |
| 6 | 0750 | -10 | 0 | 0 | -1 | 0800-0859 |
| 7 | 1452 | -8 | 0 | 0 | -1 | 1500-1559 |
| 8 | 1450 | -10 | 0 | 0 | -1 | 1500-1559 |

| | TaxiOut | WheelsOff | WheelsOn | TaxiIn | CRSArrTime | ArrTime | ArrDelay | ArrDelayMinutes |
|---|---------|-----------|----------|--------|------------|---------|----------|-----------------|
| 1 | 11.00 | 0808 | 0833 | 20.00 | 0905 | 0853 | -12 | 0 |
| 2 | 19.00 | 0814 | 0851 | 6.00 | 0905 | 0857 | -8 | 0 |
| 3 | 14.00 | 0809 | 0837 | 7.00 | 0905 | 0844 | -21 | 0 |
| 4 | 13.00 | 0807 | 0845 | 3.00 | 0905 | 0848 | -17 | 0 |
| 5 | 17.00 | 0816 | 0844 | 5.00 | 0905 | 0849 | -16 | 0 |
| 6 | 17.00 | 0807 | 0845 | 7.00 | 0905 | 0852 | -13 | 0 |
| 7 | 26.00 | 1518 | 1643 | 6.00 | 1720 | 1649 | -31 | 0 |
| 8 | 16.00 | 1506 | 1650 | 5.00 | 1720 | 1655 | -25 | 0 |

| | ArrDel15 | ArrivalDelayGroups | ArrTimeBlk | Cancelled | CancellationCode | Diverted |
|---|----------|--------------------|------------|-----------|------------------|----------|
| 1 | 0 | -1 | 0900-0959 | 0 | <NA> | 0 |
| 2 | 0 | -1 | 0900-0959 | 0 | <NA> | 0 |
| 3 | 0 | -2 | 0900-0959 | 0 | <NA> | 0 |
| 4 | 0 | -2 | 0900-0959 | 0 | <NA> | 0 |
| 5 | 0 | -2 | 0900-0959 | 0 | <NA> | 0 |
| 6 | 0 | -1 | 0900-0959 | 0 | <NA> | 0 |
| 7 | 0 | -2 | 1700-1759 | 0 | <NA> | 0 |
| 8 | 0 | -2 | 1700-1759 | 0 | <NA> | 0 |

| | CRSElapsedTime | ActualElapsedTime | AirTime | Flights | Distance | DistanceGroup |
|---|----------------|-------------------|---------|---------|----------|---------------|
| 1 | 65 | 56 | 25.00 | 1 | 101 | 1 |
| 2 | 65 | 62 | 37.00 | 1 | 101 | 1 |
| 3 | 65 | 49 | 28.00 | 1 | 101 | 1 |
| 4 | 65 | 54 | 38.00 | 1 | 101 | 1 |
| 5 | 65 | 50 | 28.00 | 1 | 101 | 1 |
| 6 | 65 | 62 | 38.00 | 1 | 101 | 1 |
| 7 | 140 | 117 | 85.00 | 1 | 585 | 3 |
| 8 | 140 | 125 | 104.00 | 1 | 585 | 3 |

| | CarrierDelay | WeatherDelay | NASDelay | SecurityDelay | LateAircraftDelay |
|---|--------------|--------------|----------|---------------|-------------------|
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |

| | | | | | |
|---|----------------------|------------------|-----------------|--------------------|----------------|
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | FirstDepTime | TotalAddGTime | LongestAddGTime | DivAirportLandings | DivReachedDest |
| 1 | <NA> | <NA> | <NA> | 0 | <NA> |
| 2 | <NA> | <NA> | <NA> | 0 | <NA> |
| 3 | <NA> | <NA> | <NA> | 0 | <NA> |
| 4 | <NA> | <NA> | <NA> | 0 | <NA> |
| 5 | <NA> | <NA> | <NA> | 0 | <NA> |
| 6 | <NA> | <NA> | <NA> | 0 | <NA> |
| 7 | <NA> | <NA> | <NA> | 0 | <NA> |
| 8 | <NA> | <NA> | <NA> | 0 | <NA> |
| | DivActualElapsedTime | DivArrDelay | DivDistance | Div1Airport | Div1AirportID |
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | Div1AirportSeqID | Div1WheelsOn | Div1TotalGTime | Div1LongestGTime | Div1WheelsOff |
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | Div1TailNum | Div2Airport | Div2AirportID | Div2AirportSeqID | Div2WheelsOn |
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | Div2TotalGTime | Div2LongestGTime | Div2WheelsOff | Div2TailNum | Div3Airport |
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |

| | | | | | |
|---|------------------|------------------|------------------|----------------|------------------|
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | Div3AirportID | Div3AirportSeqID | Div3WheelsOn | Div3TotalGTime | Div3LongestGTime |
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | Div3WheelsOff | Div3TailNum | Div4Airport | Div4AirportID | Div4AirportSeqID |
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | Div4WheelsOn | Div4TotalGTime | Div4LongestGTime | Div4WheelsOff | Div4TailNum |
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | Div5Airport | Div5AirportID | Div5AirportSeqID | Div5WheelsOn | Div5TotalGTime |
| 1 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 2 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 3 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 4 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> | <NA> |
| | Div5LongestGTime | Div5WheelsOff | Div5TailNum | column109 | |
| 1 | <NA> | <NA> | <NA> | <NA> | |
| 2 | <NA> | <NA> | <NA> | <NA> | |
| 3 | <NA> | <NA> | <NA> | <NA> | |

| | | | | |
|---|------|------|------|------|
| 4 | <NA> | <NA> | <NA> | <NA> |
| 5 | <NA> | <NA> | <NA> | <NA> |
| 6 | <NA> | <NA> | <NA> | <NA> |
| 7 | <NA> | <NA> | <NA> | <NA> |
| 8 | <NA> | <NA> | <NA> | <NA> |

1. SQL queries via the DBI package

- What is the average Departure Delay (and number of flights) for each destination?
- Can now apply the `MEDIAN()` function!

```
```{r}
DBI::dbGetQuery(con_parq,
 "SELECT MEDIAN(DepDelay) AS med_delay,
 AVG(DepDelay) AS mean_delay, Dest
 FROM read_parquet('activity/data_airlines/Year*/*.parquet')
 GROUP BY Dest
 LIMIT 8;")
```
```

| | med_delay | mean_delay | Dest |
|---|-----------|------------|------|
| 1 | -1 | 11.344920 | STL |
| 2 | -2 | 11.778499 | CHS |
| 3 | -3 | 11.035951 | GRB |
| 4 | -2 | 12.511458 | SYR |
| 5 | -2 | 11.553647 | ROC |
| 6 | -2 | 11.161244 | CMH |
| 7 | -3 | 10.076172 | MDT |
| 8 | -4 | 5.933875 | MLU |

2. Translating dplyr code into SQL

- Creating a `tbl` requires a pointer to the folder where the parquet files live.

```
```{r}
flights_parq <- tbl(
 con_parq,
 "read_parquet('activity/data_airlines/Year*/*.parquet')")
flights_parq |>
```

```

head()
...

Source: SQL [?? x 110]
Database: DuckDB v1.1.2 [root@Darwin 24.5.0:R 4.4.2/:memory:]
 Year Quarter Month DayOfMonth DayOfWeek FlightDate
 <dbl> <dbl> <dbl> <dbl> <dbl> <date>
1 2023 1 1 2 1 2023-01-02
2 2023 1 1 3 2 2023-01-03
3 2023 1 1 4 3 2023-01-04
4 2023 1 1 5 4 2023-01-05
5 2023 1 1 6 5 2023-01-06
6 2023 1 1 7 6 2023-01-07
i 104 more variables: Reporting_Airline <chr>,
DOT_ID_Reporting_Airline <dbl>,
IATA_CODE_Reporting_Airline <chr>, Tail_Number <chr>,
Flight_Number_Reporting_Airline <dbl>, OriginAirportID <dbl>,
OriginAirportSeqID <dbl>, OriginCityMarketID <dbl>, Origin <chr>,
OriginCityName <chr>, OriginState <chr>, OriginStateFips <chr>,
OriginStateName <chr>, OriginWac <dbl>, DestAirportID <dbl>, ...

```

## 2. Translating dplyr code into SQL

- Over what years is the `flights` data taken?

```

```{r}
yrs <- flights_parq |>
  summarize(min_year = min(Year, na.rm = TRUE),
            max_year = max(Year, na.rm = TRUE))

yrs

dplyr::show_query(yrs)
...

```

```

# Source:   SQL [?? x 2]
# Database: DuckDB v1.1.2 [root@Darwin 24.5.0:R 4.4.2/:memory:]
  min_year max_year
    <dbl>   <dbl>
1    2023    2024
<SQL>

```



```
SELECT MIN("Year") AS min_year, MAX("Year") AS max_year
FROM (FROM read_parquet('activity/data_airlines/Year*/*.parquet')) q01
```

2. Translating dplyr code into SQL

- Create a data set containing only flights between LAX and BOS in 2024.

```
```{r}
#| cache: true

la_bos_parquet <- flights_parq |>
 filter(Year == 2024 & ((Origin == "LAX" & Dest == "BOS") |
 (Origin == "BOS" & Dest == "LAX")))

dplyr::show_query(la_bos_parquet)
```
```

```
<SQL>
SELECT q01.*
FROM (FROM read_parquet('activity/data_airlines/Year*/*.parquet')) q01
WHERE ("Year" = 2024.0 AND ((Origin = 'LAX' AND Dest = 'BOS') OR (Origin = 'BOS' AND Dest =
```

2. Translating dplyr code into SQL

- What is the average Departure Delay (and number of flights) for each destination?
- Can now apply the MEDIAN() function!

```
```{r}
#| cache: true

aveDepDel_parquet <- flights_parq |>
 group_by(Dest) |>
 summarize(mean_delay = mean(DepDelay),
 med_delay = median(DepDelay),
 num_flights = n())

aveDepDel_parquet

dplyr::show_query(aveDepDel_parquet)
```
```


[illegible]

3. Direct SQL queries via `sql` chunks

- What is the average Departure Delay (and number of flights) for each destination?
- Can now apply the `MEDIAN()` function!

```

```{sql}
#| connection: con_parq
#| cache: true

SELECT MEDIAN(DepDelay) AS med_delay,
 AVG(DepDelay) AS mean_delay, Dest
FROM read_parquet('activity/data_airlines/Year*/*.parquet')
GROUP BY Dest
LIMIT 8;
```

```

Table 5: 8 records

| med_delay | mean_delay | Dest |
|-----------|------------|------|
| -2 | 9.750851 | PDX |
| -1 | 13.586533 | MSY |
| -2 | 16.059348 | BOS |
| -2 | 8.899172 | SNA |
| -1 | 13.173580 | ALB |
| -1 | 13.369200 | PVD |
| -1 | 10.840029 | SAN |
| -4 | 3.948148 | KTN |

Good practice

Always a good idea to terminate the **SQL** connection when you are done with it.

```
RMariaDB::dbDisconnect(con_parq, shutdown = TRUE)
```